Model Interpretability and Data Drift

Nicholas Moore

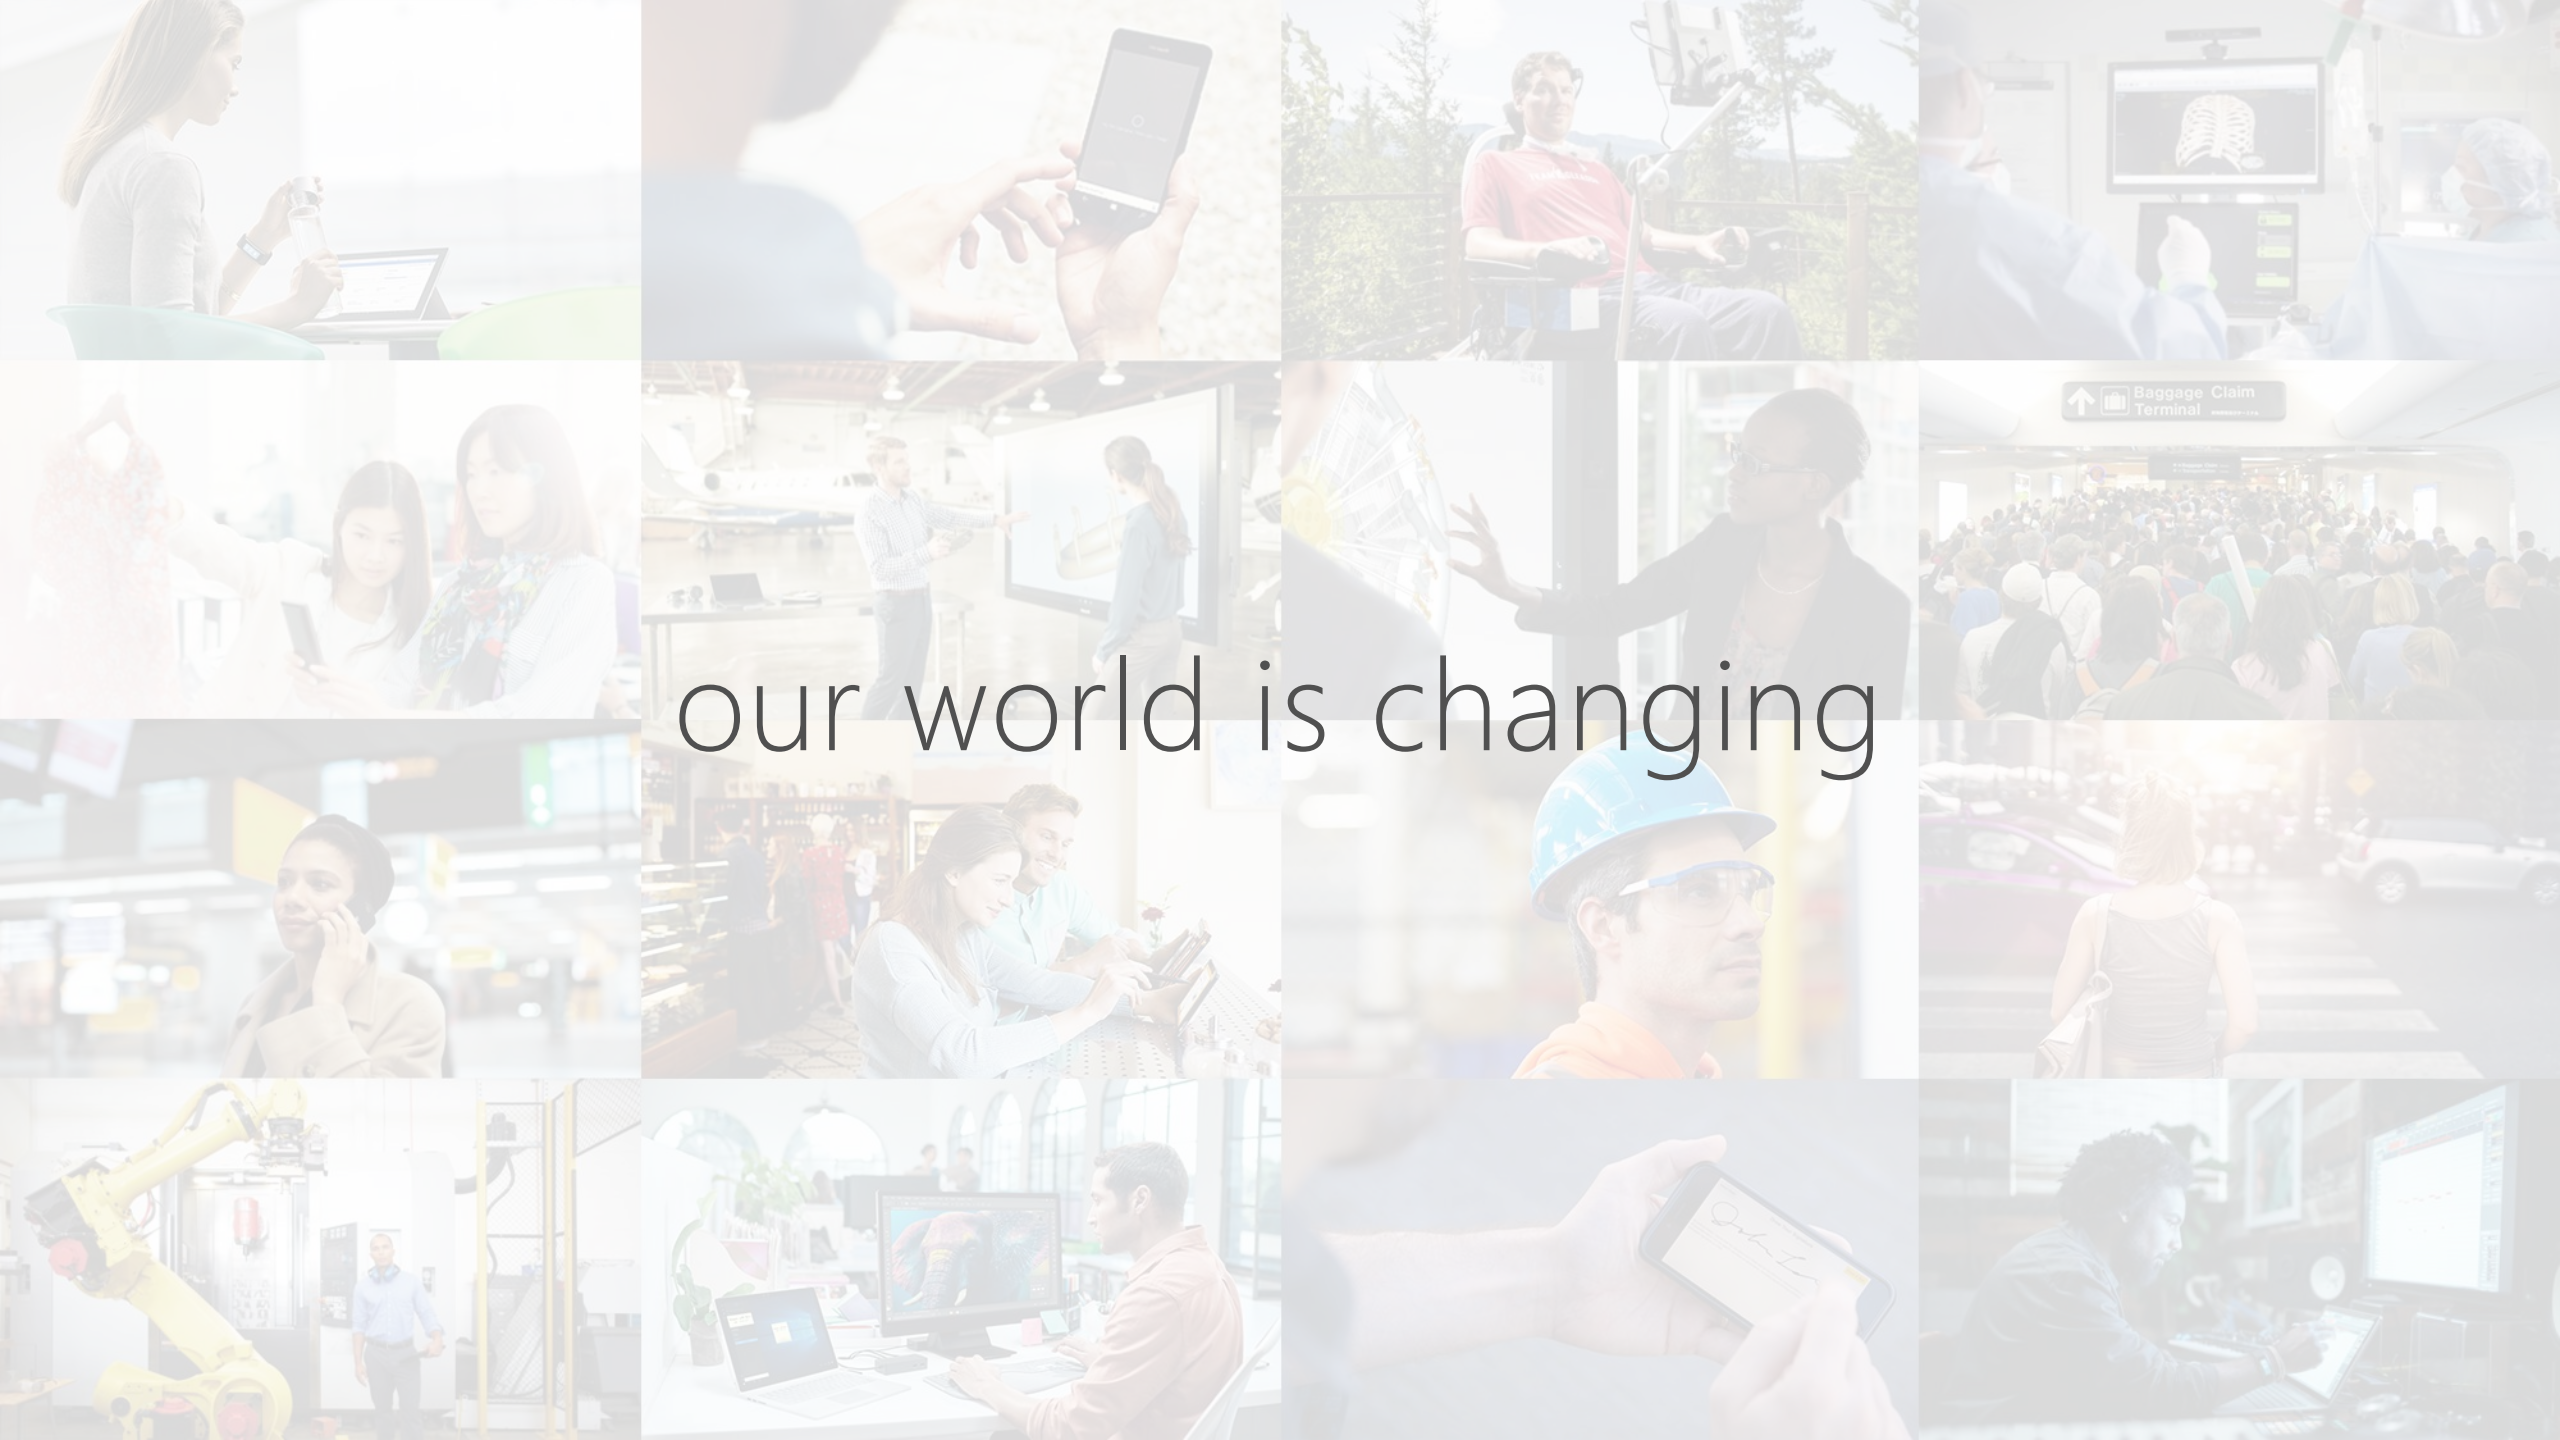

our world is changing

Data Science Lifecycle

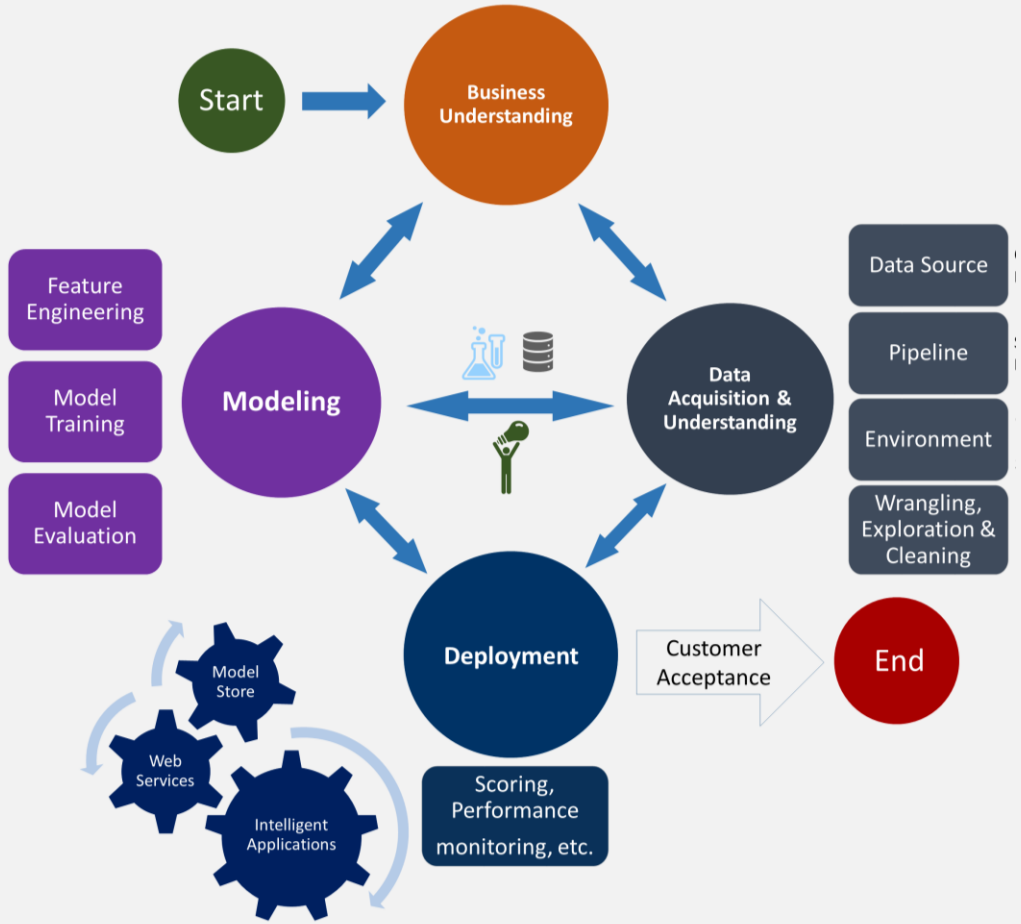

Team Data Science Process: https://docs.microsoft.com/en-us/azure/machine-learning/team-data-science-process/overview

there is a lot of hype around model creation, but not model maintenance

The moment you put a model in production, it starts degrading.

Covariate Shift

Changes in the distribution of independent variables

Concept Drift

Changes in the relationship between independent and target variables

Model Interpretability

Understanding the process a model uses to arrive at an outcome

MACHINE LEARNING BASICS

SUPERVISED MACHINE LEARNING

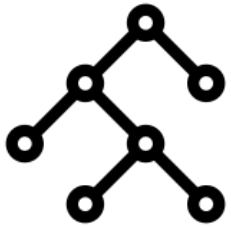

Classification

Categorical outputs

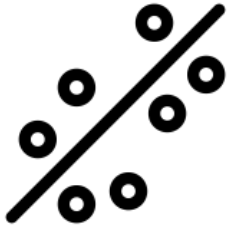

Regression

Continuous outputs

Reproduce outputs from a training data set by creating a rule that maps inputs to outputs

MODEL CREATION PROCESS

CLASSIFICATION

Iterativly find the best boundry
to separate classes

Use boundry to classify new
data into a class

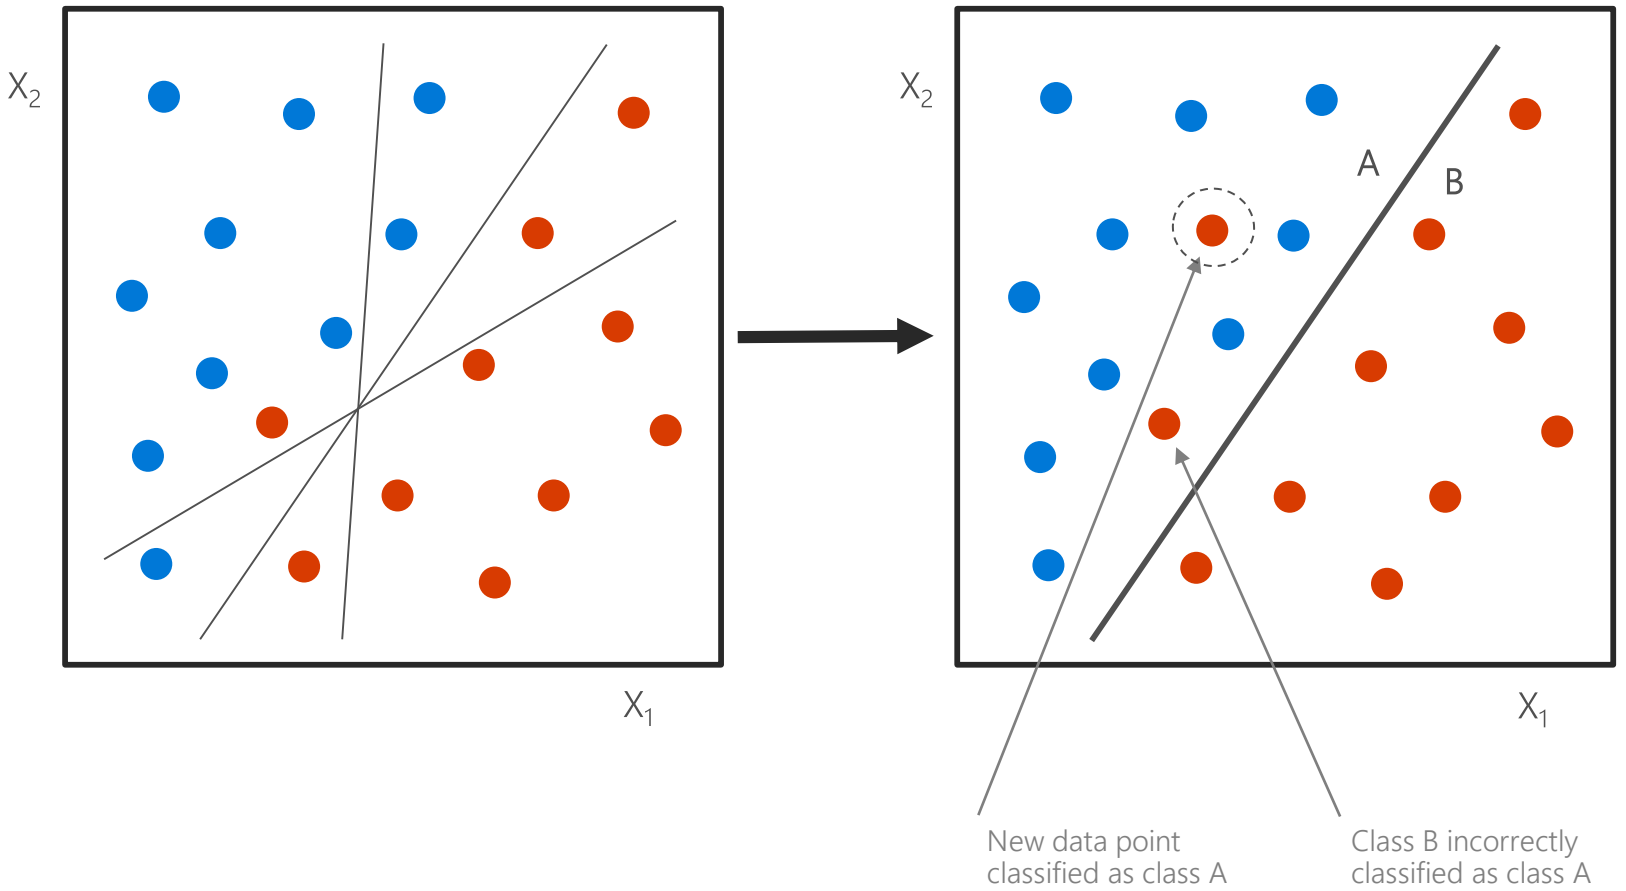

MODEL CREATION PROCESS

REGRESSION

Iterativly find the line of best fit to minimise the distance between the target value

Use the line of best fit to predict the target value

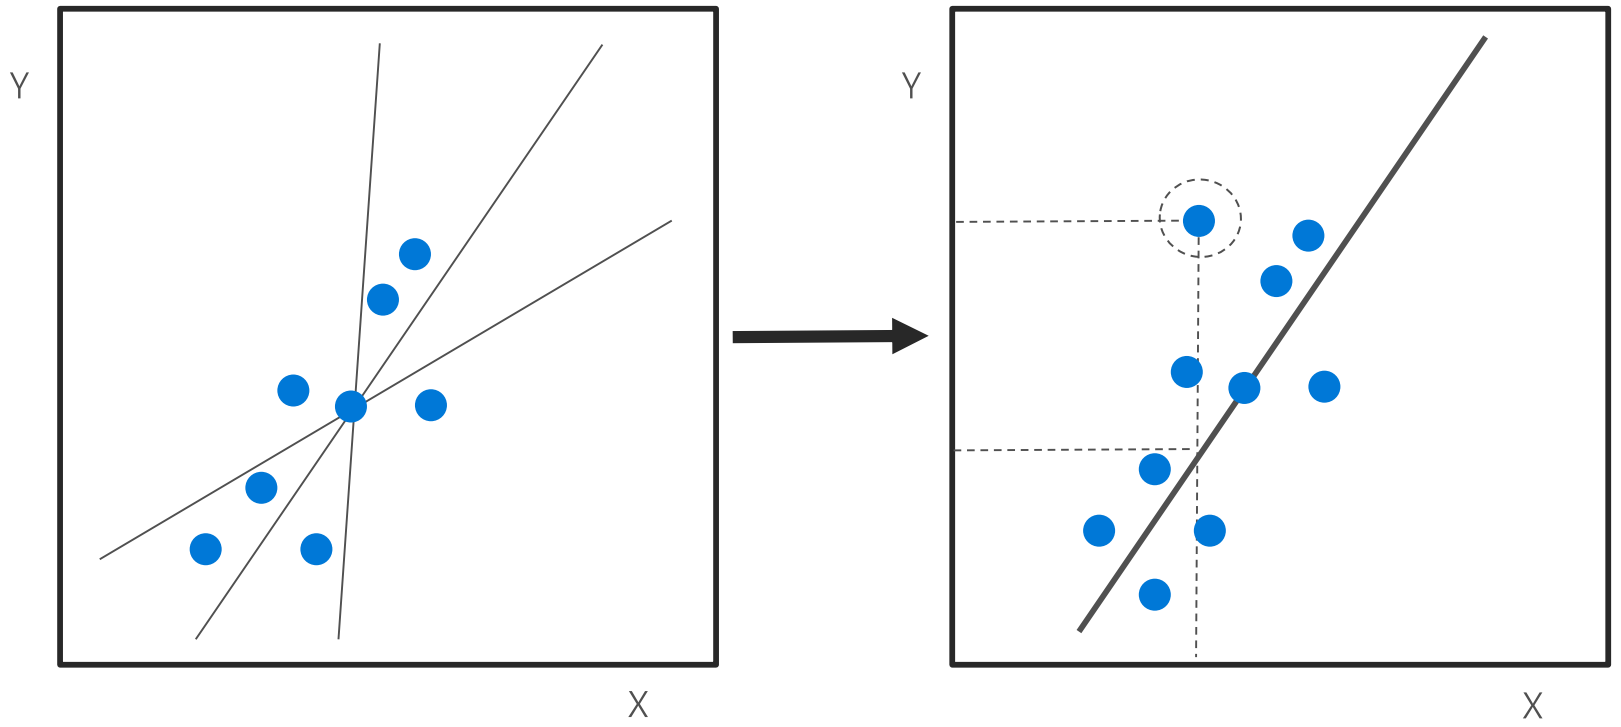

DATA DRIFT

CONCEPT DRIFT

CLASSIFICATION

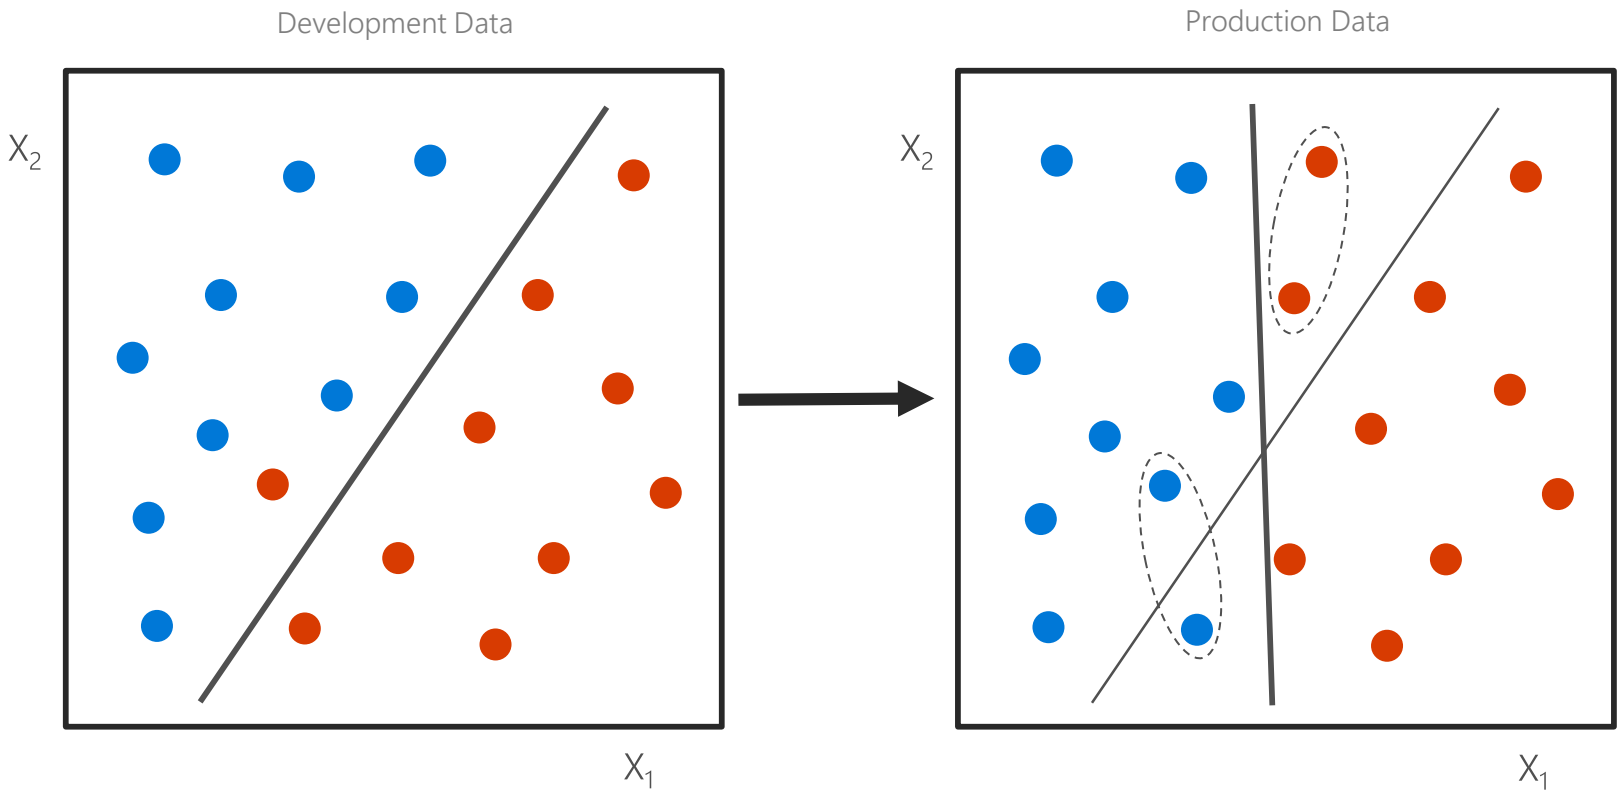

Change in the relationship between independent
and target variables in the underlying problem overtime

CONCEPT DRIFT

REGRESSION

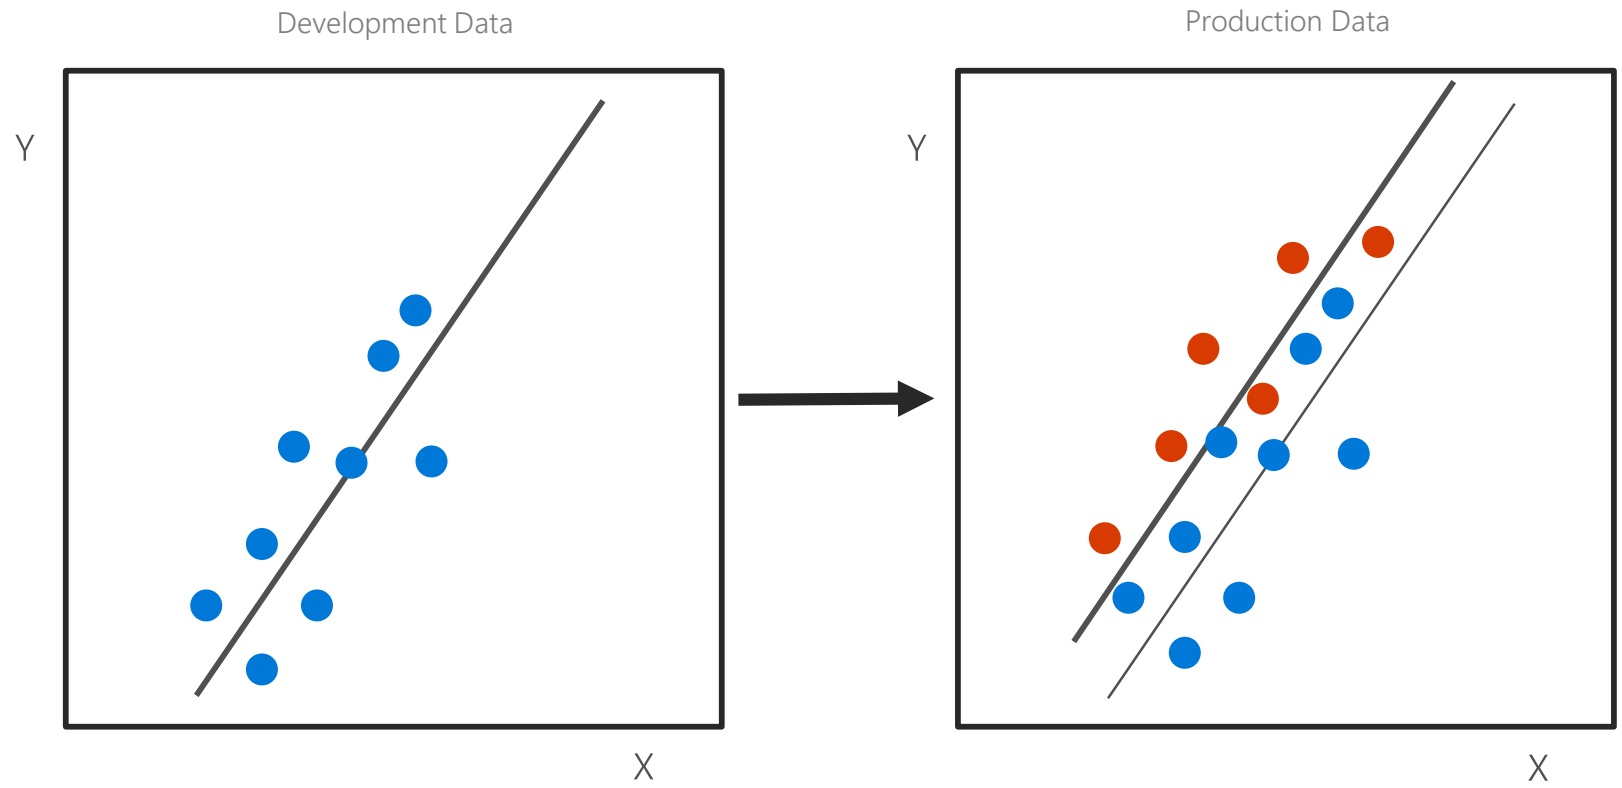

Change in the relationship between independent
and target variables in the underlying problem overtime

COVARIATE SHIFT

CLASSIFICATION

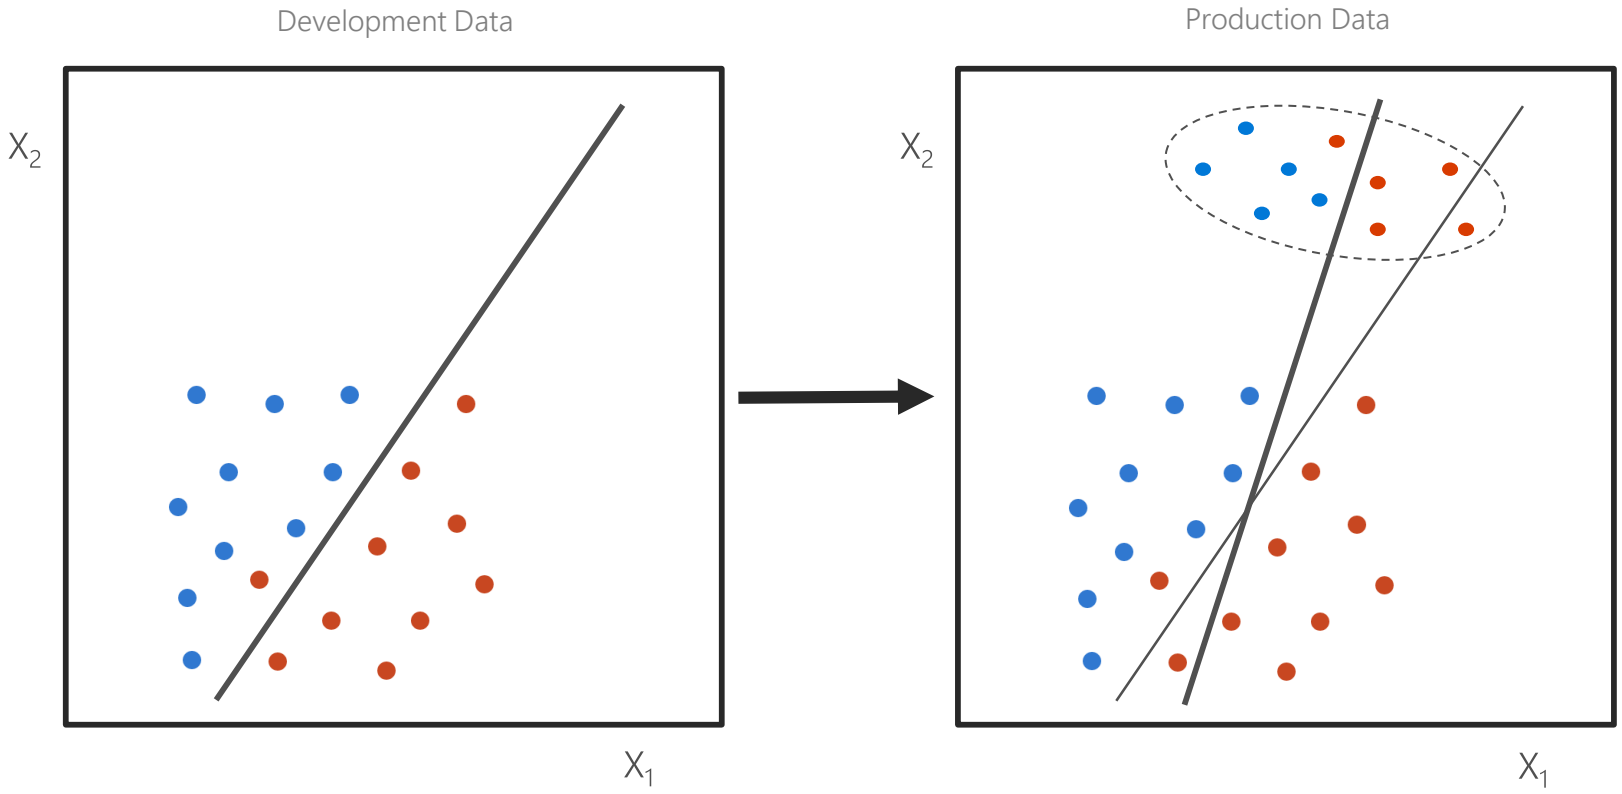

Change in the distribution of the independent variables in the underlying problem overtime

COVARIATE SHIFT

REGRESSION

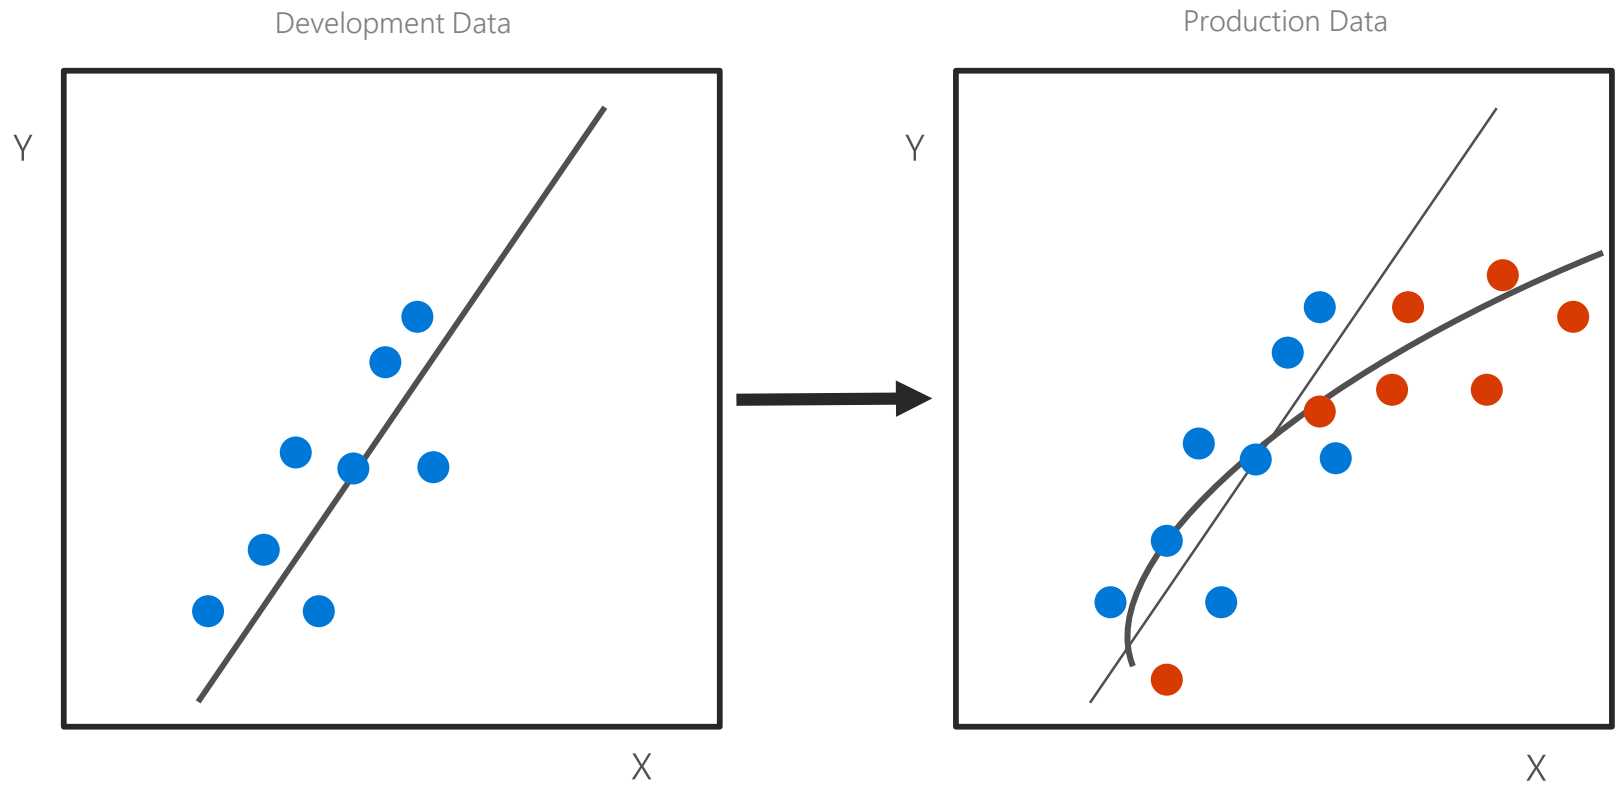

Change in the distribution of the independent variables in the underlying problem overtime

Why is this a problem?

Production model no longer fit-for-purpose

Can be difficult to detect

Requires models to be monitored in production

Changes can be gradual, cyclical, or abrupt

Increased model maintenance

How to address the problem?

- Periodically re-fit or update the model
- Monitor distribution of independent variables in production
- Weight data
- Periodically assess the performance of your predictions
- Learn the change

EXAMPLES OF MODEL DRIFT

CONCEPT DRIFT

Identifying patterns of people who commit fraud or hacking computer networks

New regulations introduced to save electricity which influences predicting demand

New production procedures which impact models designed to assess quality of products / produce

COVARIATE SHIFT

Changes in common words or meaning of works overtime

Ware and tear of sensors over time in predictive maintenance solutions

Different accents in speech-to-text solutions

MODEL INTERPRETABILITY

Fair and Accountable ML

What outcomes do machine learning models influence:

Will a person get a loan or credit card from a lending institution?

Will a teacher be fired based on their teaching evaluation score?

Will a prisoner be released on parole based on likelihood to reoffend?

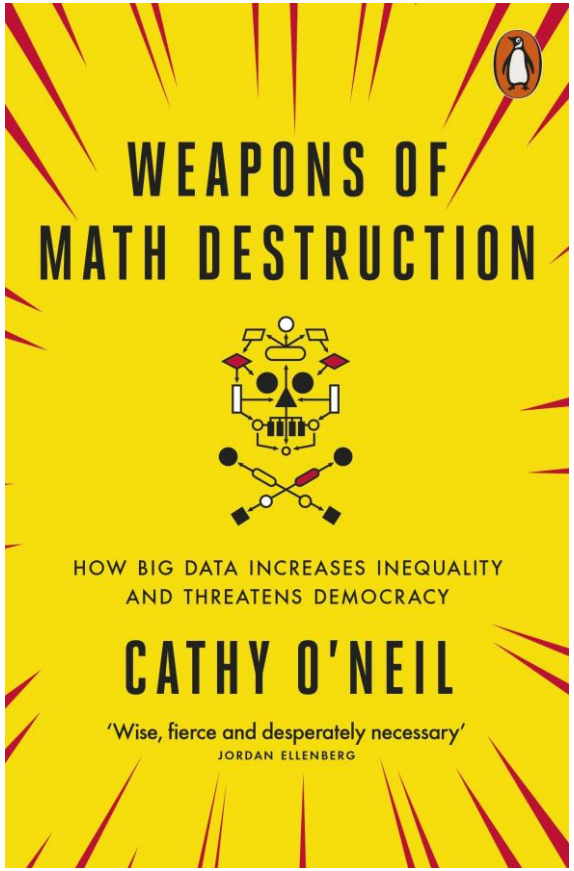

Why is it important?

Model interpratability aims to understand the process a machine learning model uses to make predictions

- Feature causality
- Inform data collection and feature engineering
- Model debugging
- Building trust with business stakeholders
- Regulation and auditability

INTERPRETABILITY OF ALGORITHMS

WHITE-BOX MODELS

Simpler computation
Less predictive capability
Easier to understand

Linear Regression
Logistic Regression
Decision Trees
Naïve Bayes
K-Nearest Neighbors

BLACK-BOX MODELS

High computational complexity
Emphasis on predictive capability
Difficult to understand

Ensembles
Kernel-based SVM
Neural Networks
Deep Learning

Model Interpretability Techniques

PERMUTATION IMPORTANCE

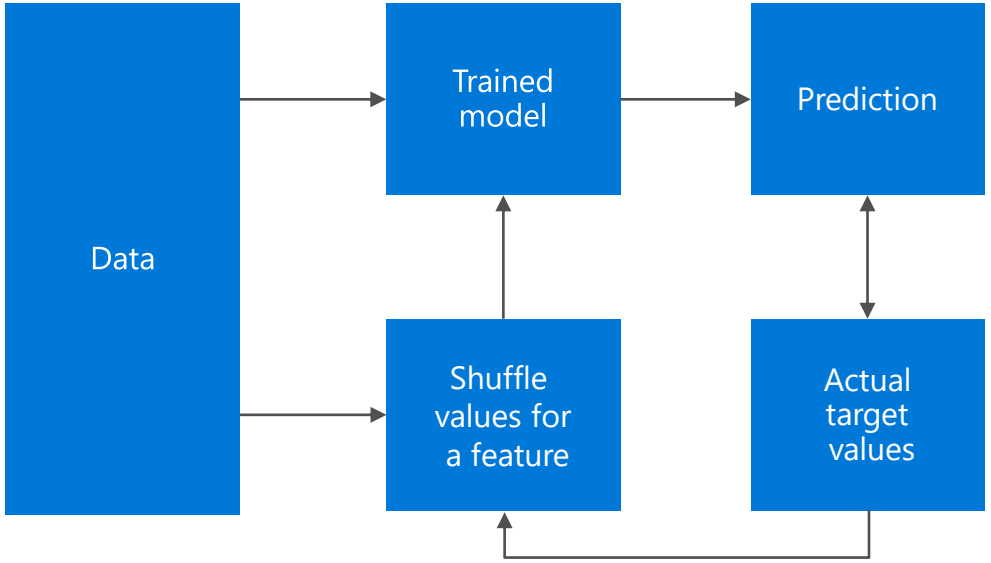

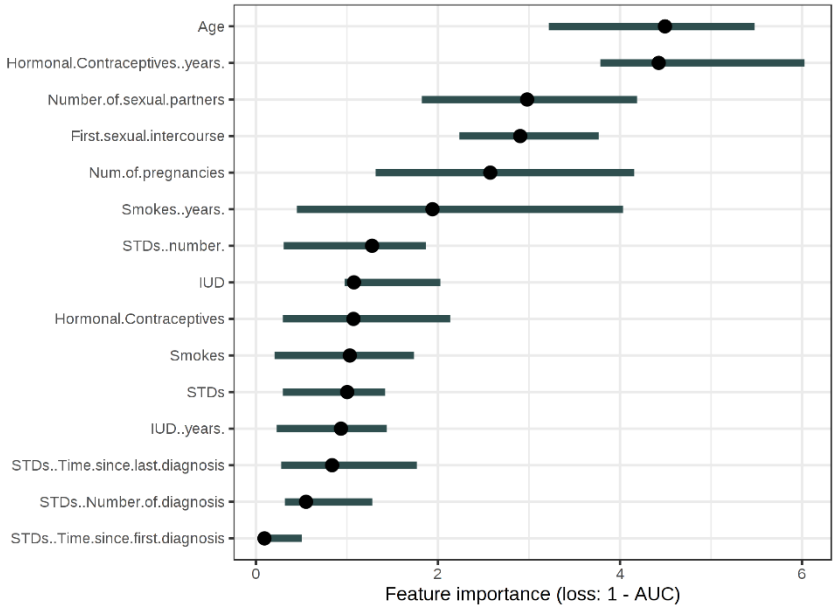

Model Interpretability Techniques

PARTIAL DEPENDENCE PLOT

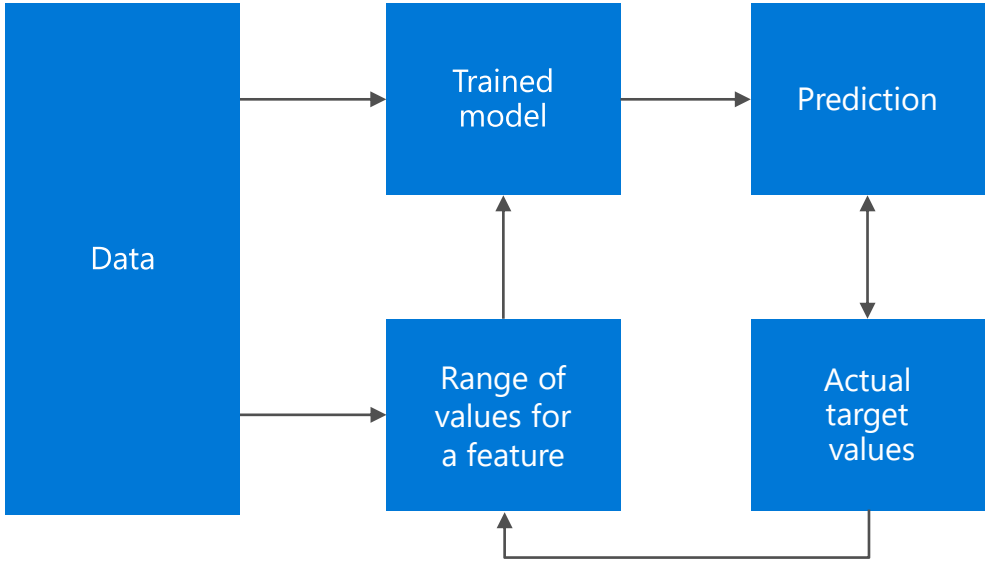

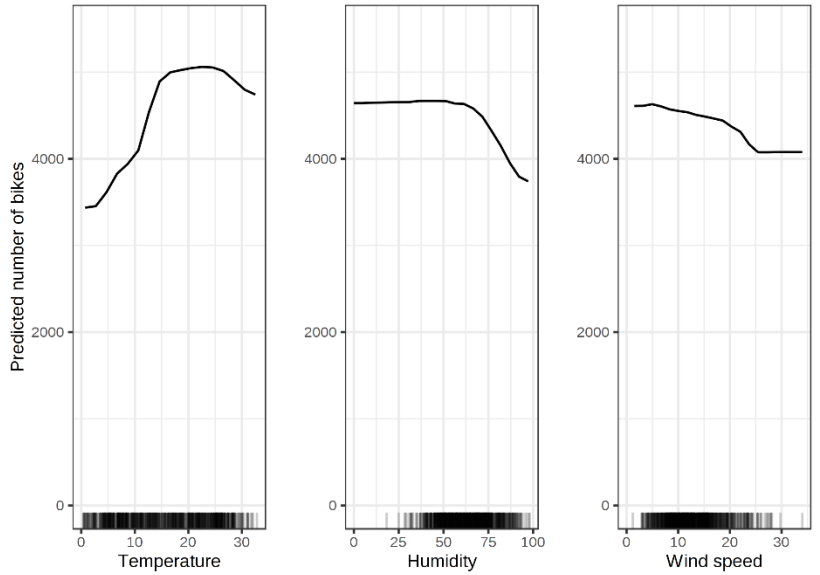

Model Interpretability Techniques

GLOBAL SURROGATE MODELS

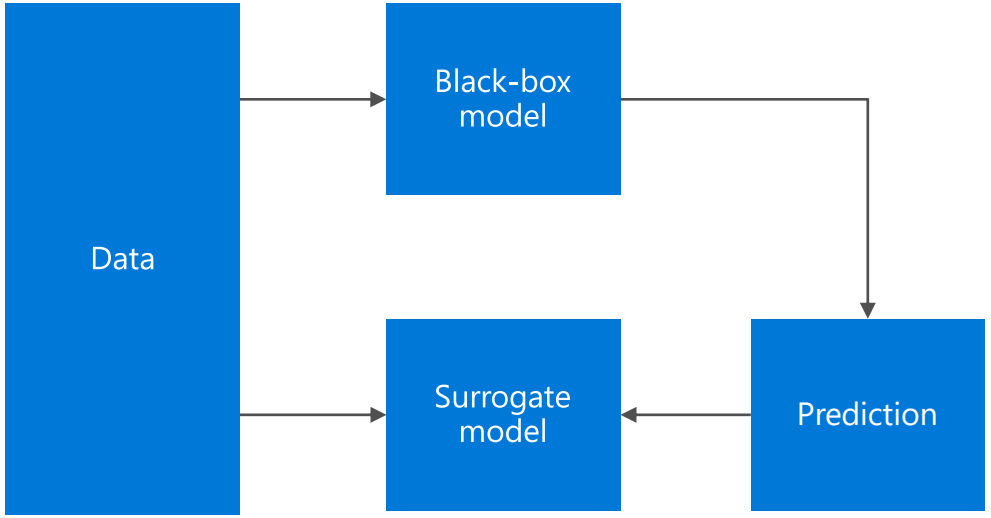

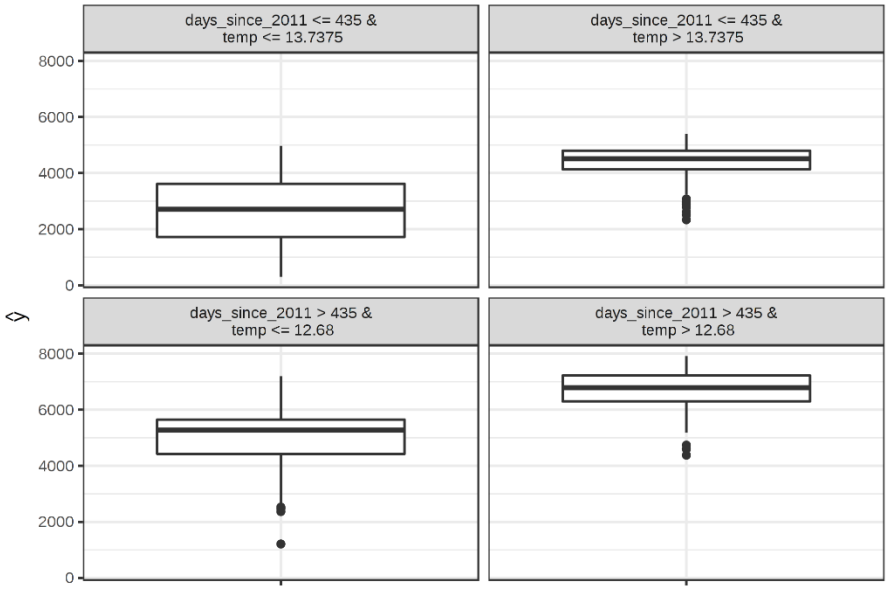

Model Interpretability Techniques

LOCAL INTERPRETABLE MODEL AGNOSTIC EXPECTATIONS (LIME)

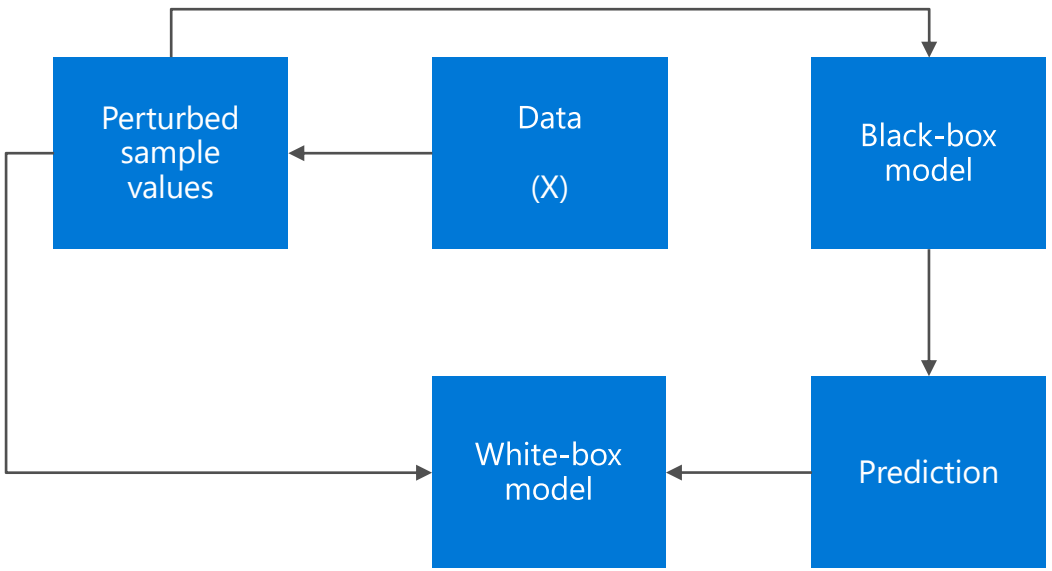

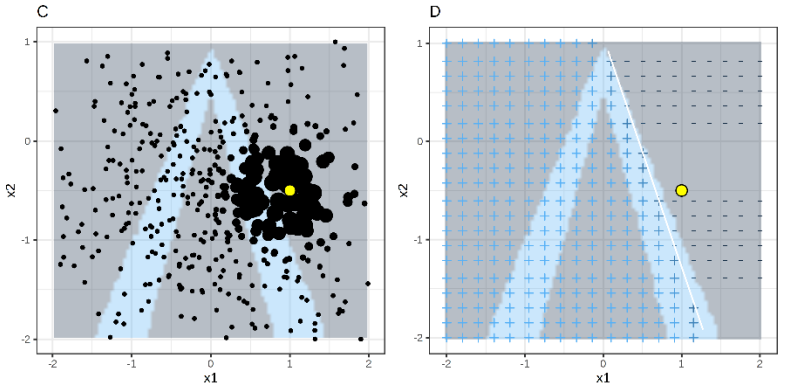

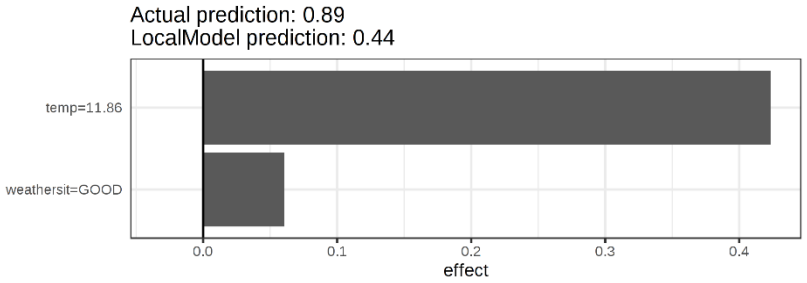

Model Interpretability Techniques

SHAPLEY ADDITIVE EXPLANATIONS (SHAP)

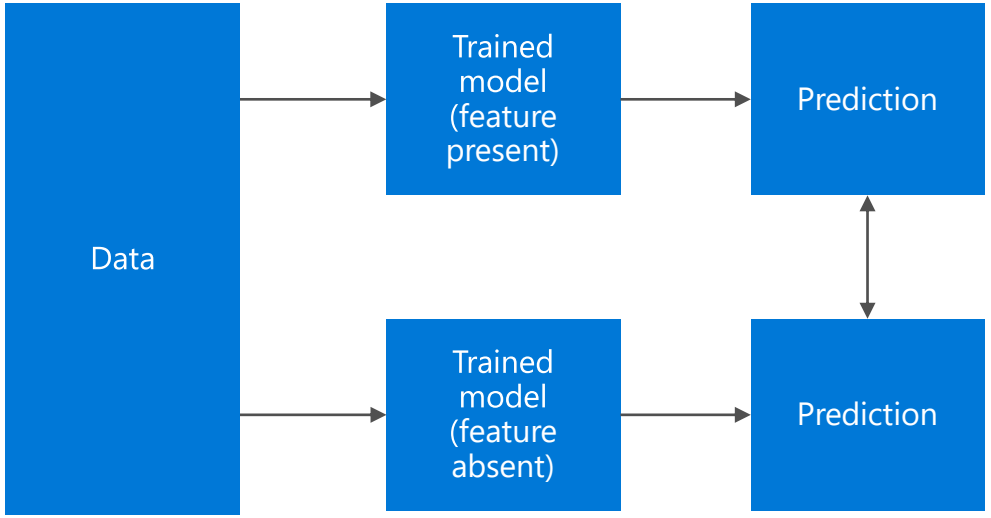

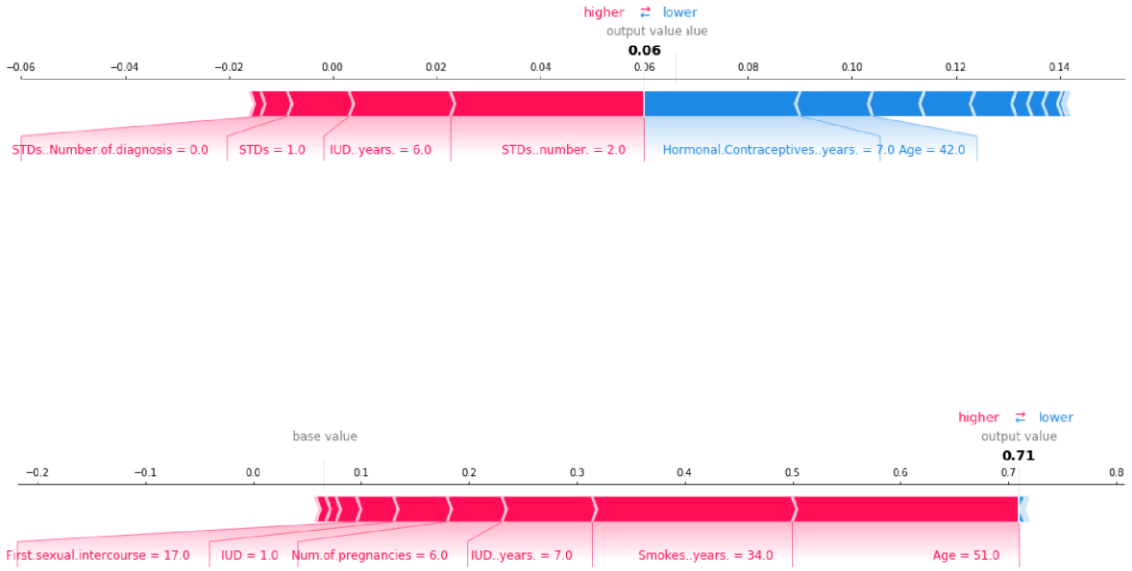

SOLUTIONS FROM MICROSOFT AZURE

Azure Machine Learning service

Set of Azure Cloud Services

Python SDK

That enables you to:

- ✓ Prepare Data
- ✓ Build Models
- ✓ Train Models
- ✓ Manage Models
- ✓ Track Experiments
- ✓ Deploy Models

Azure Machine Learning Service

DATA DRIFT

- Measures the magnitude of data drift, called the drift coefficient.
- Measures the data drift contribution by feature, informing which features caused data drift
- Measures distance metrics - currently Wasserstein and Energy Distance are computed
- Measures distributions of features - currently kernel density estimation and histograms. Send alerts to data drift by email

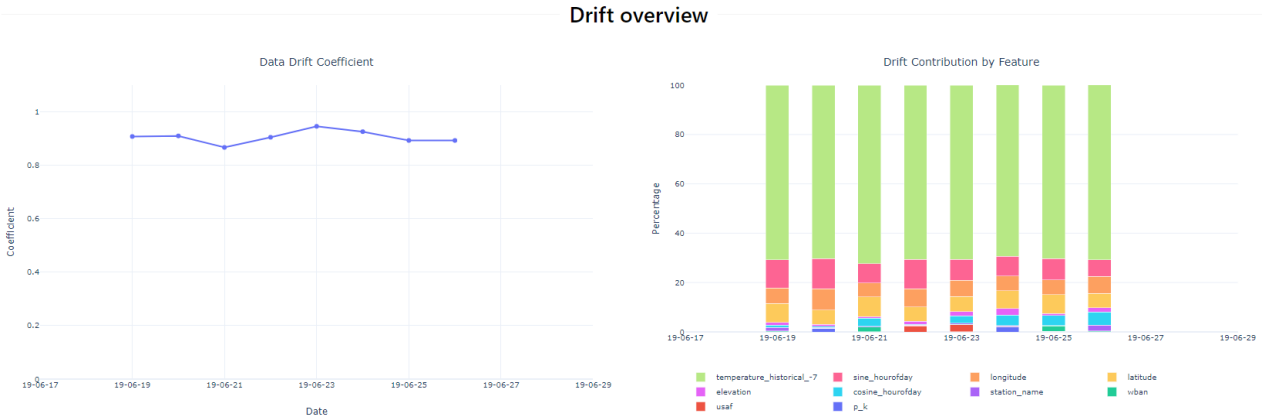

Azure Machine Learning Service

MODEL INTEPRETABILITY

- Feature importance values for both raw and engineered features
- Interpretability on real-world datasets at scale, during training and inference
- Interactive visualizations to aid you in the discovery of patterns in data and explanations at training time
- Explain machine learning models globally on all data, or locally on a specific data point using the state-of-art technologies in an easy-to-use and scalable fashion